

# Lecture 2:

# Linear classifier

Dmitry Yashunin  
IntelliVision

# Recall from last time: Challenges of recognition

Illumination



Deformation



Occlusion



Background Clutter

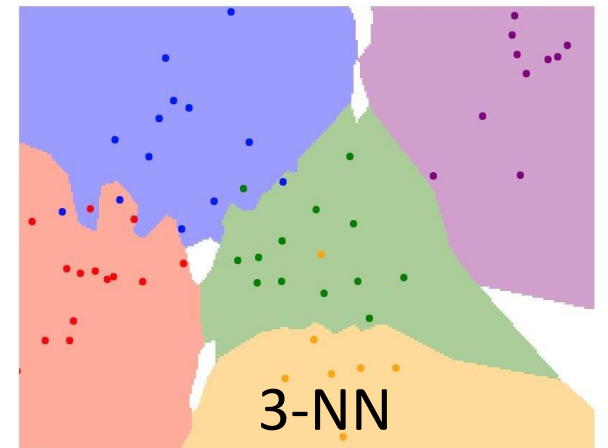
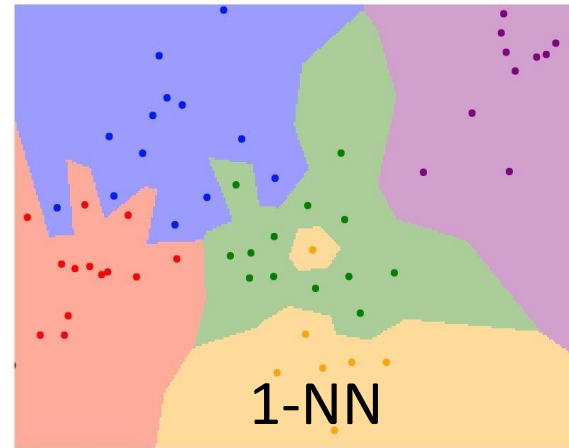


Intraclass variations



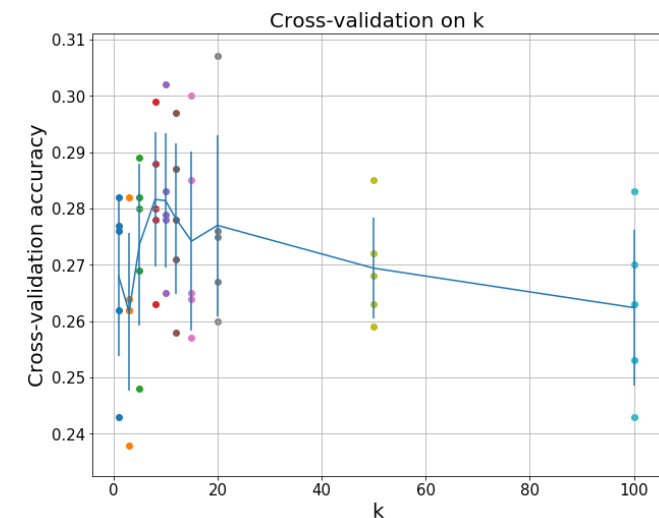
# Recall from last time: data-driven approach, kNN

CIFAR10

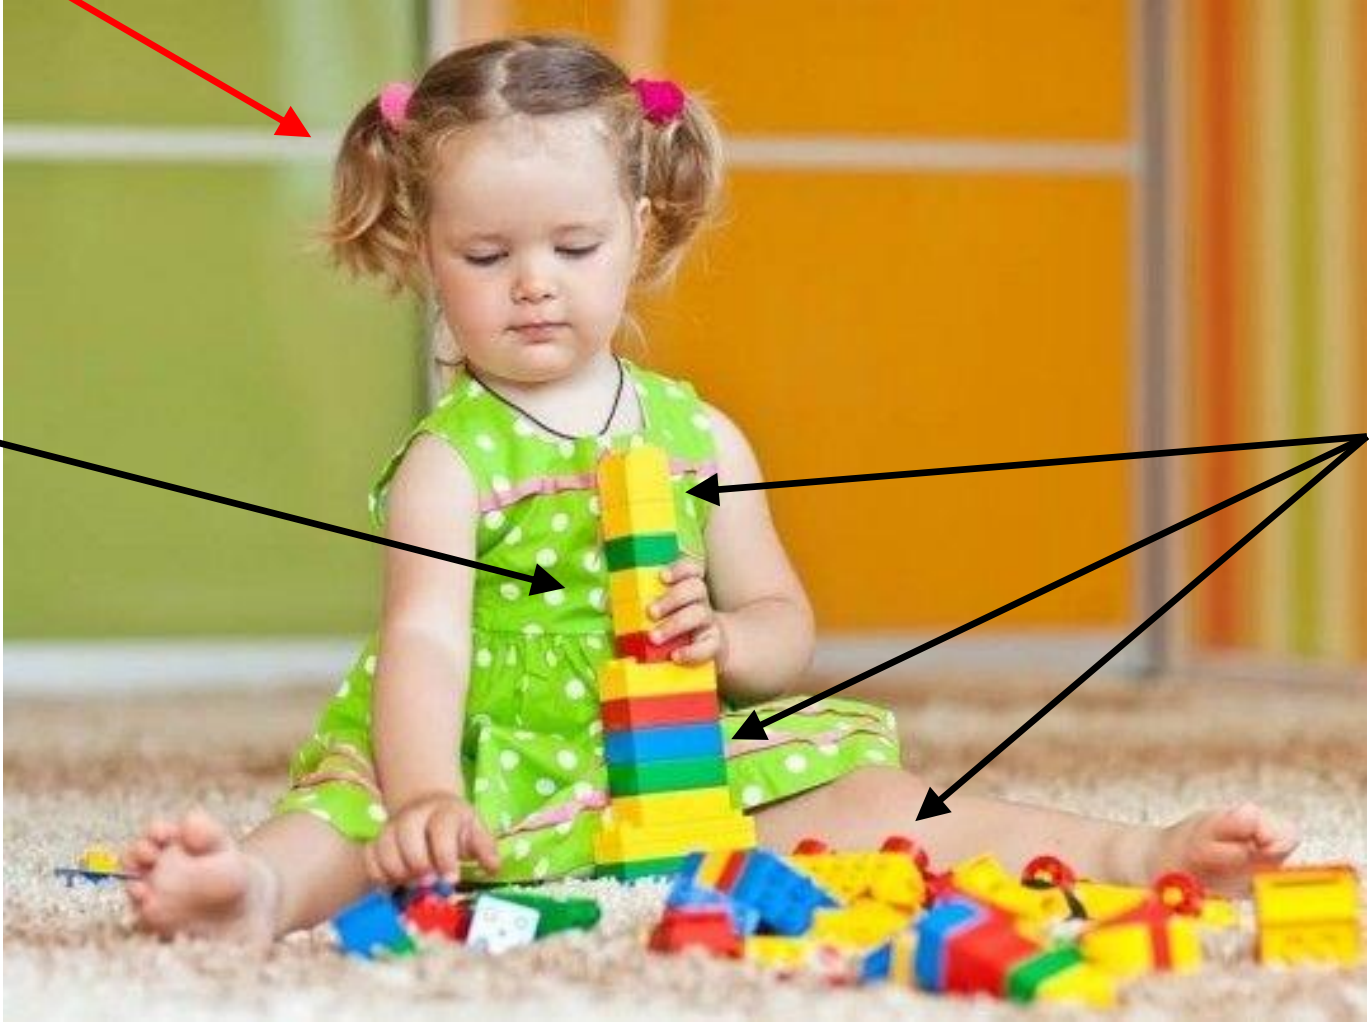


Cross-validation

fold 1	fold 2	fold 3	fold 4	test
fold 1	fold 2	fold 3	fold 4	test



Neural Network practitioner

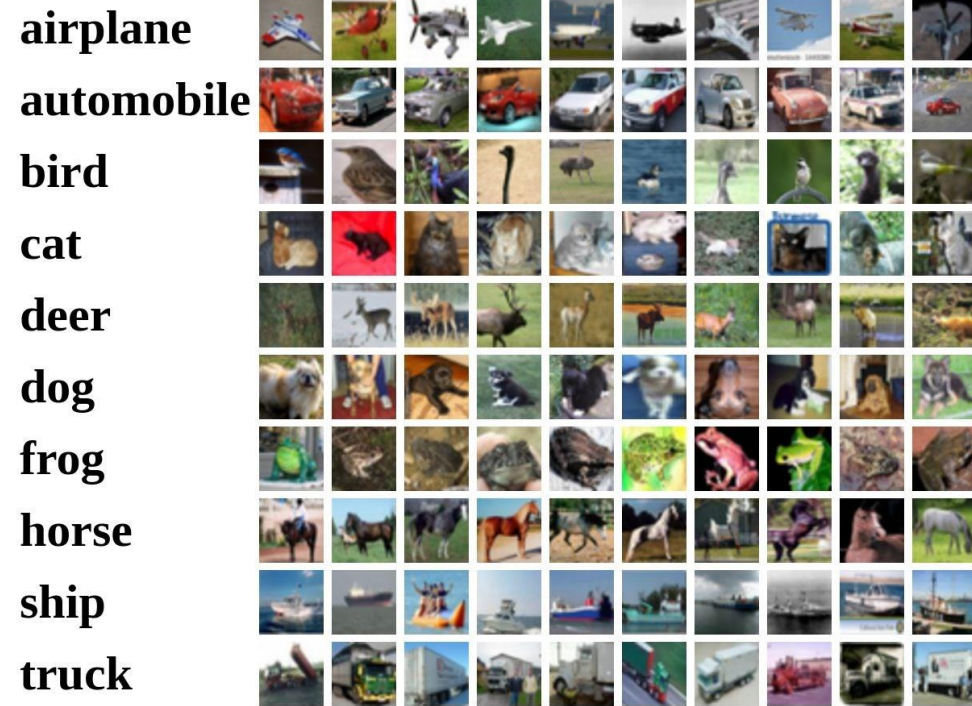


Neural Network

Linear classifiers



# Recall: CIFAR10



**10** classes

**50,000** training images

**10,000** testing images

Image size: **32x32x3**

width

height

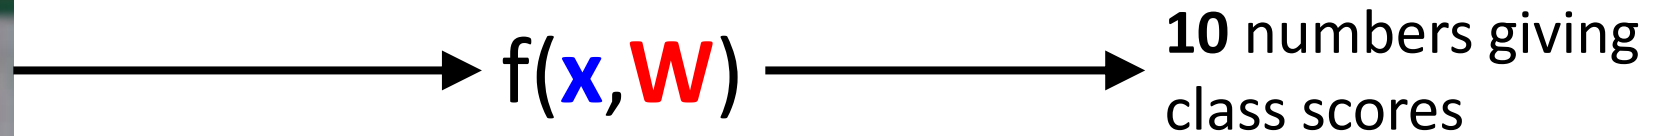
channels (rgb)

# Parametric Approach

Image



Array of **32x32x3** numbers  
(3072 numbers total)



$f(\mathbf{x}, \mathbf{W})$

10 numbers giving  
class scores



$\mathbf{W}$

parameters  
or weights

# Parametric Approach: Linear Classifier

Image



Array of **32x32x3** numbers  
(3072 numbers total)

$$f(x, W) = Wx$$

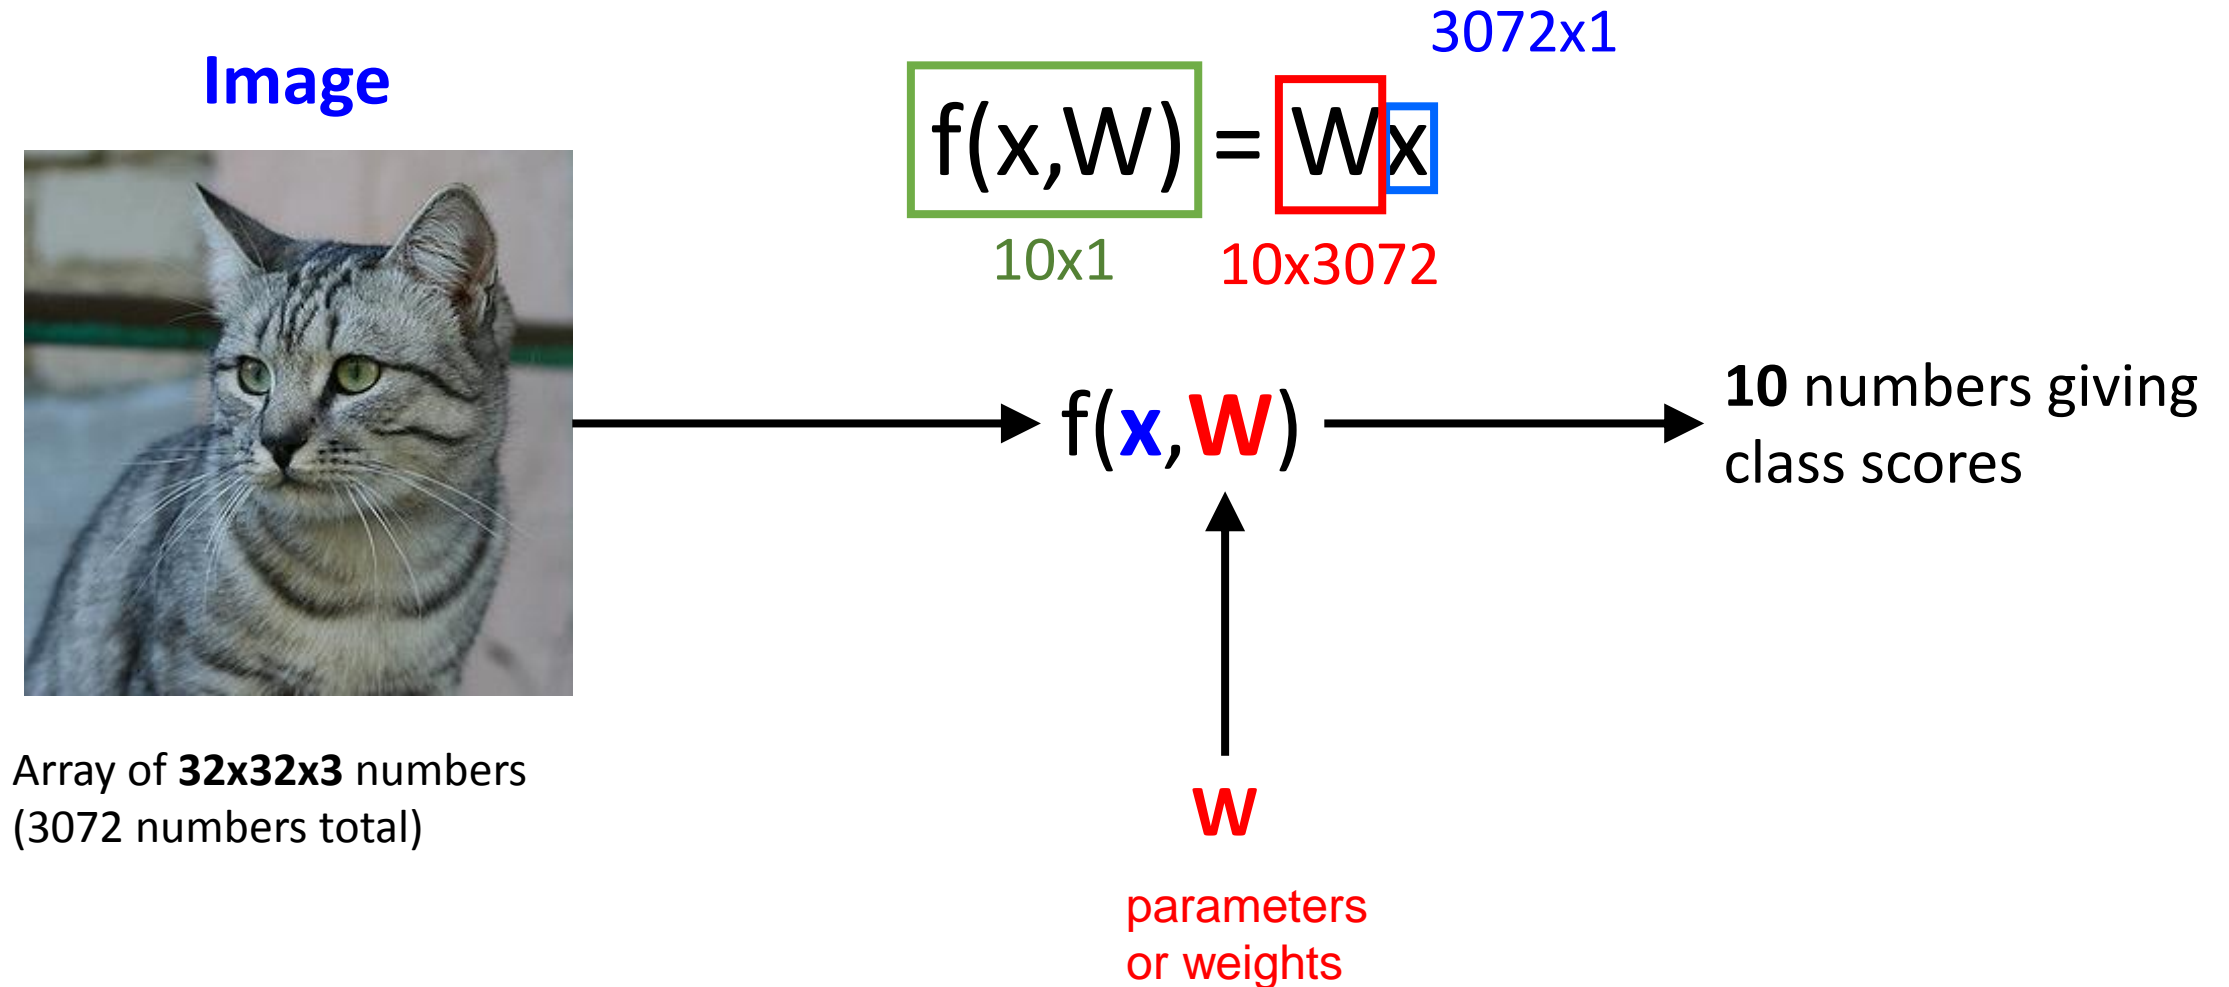
$f(\mathbf{x}, \mathbf{W})$

**10** numbers giving  
class scores

**W**

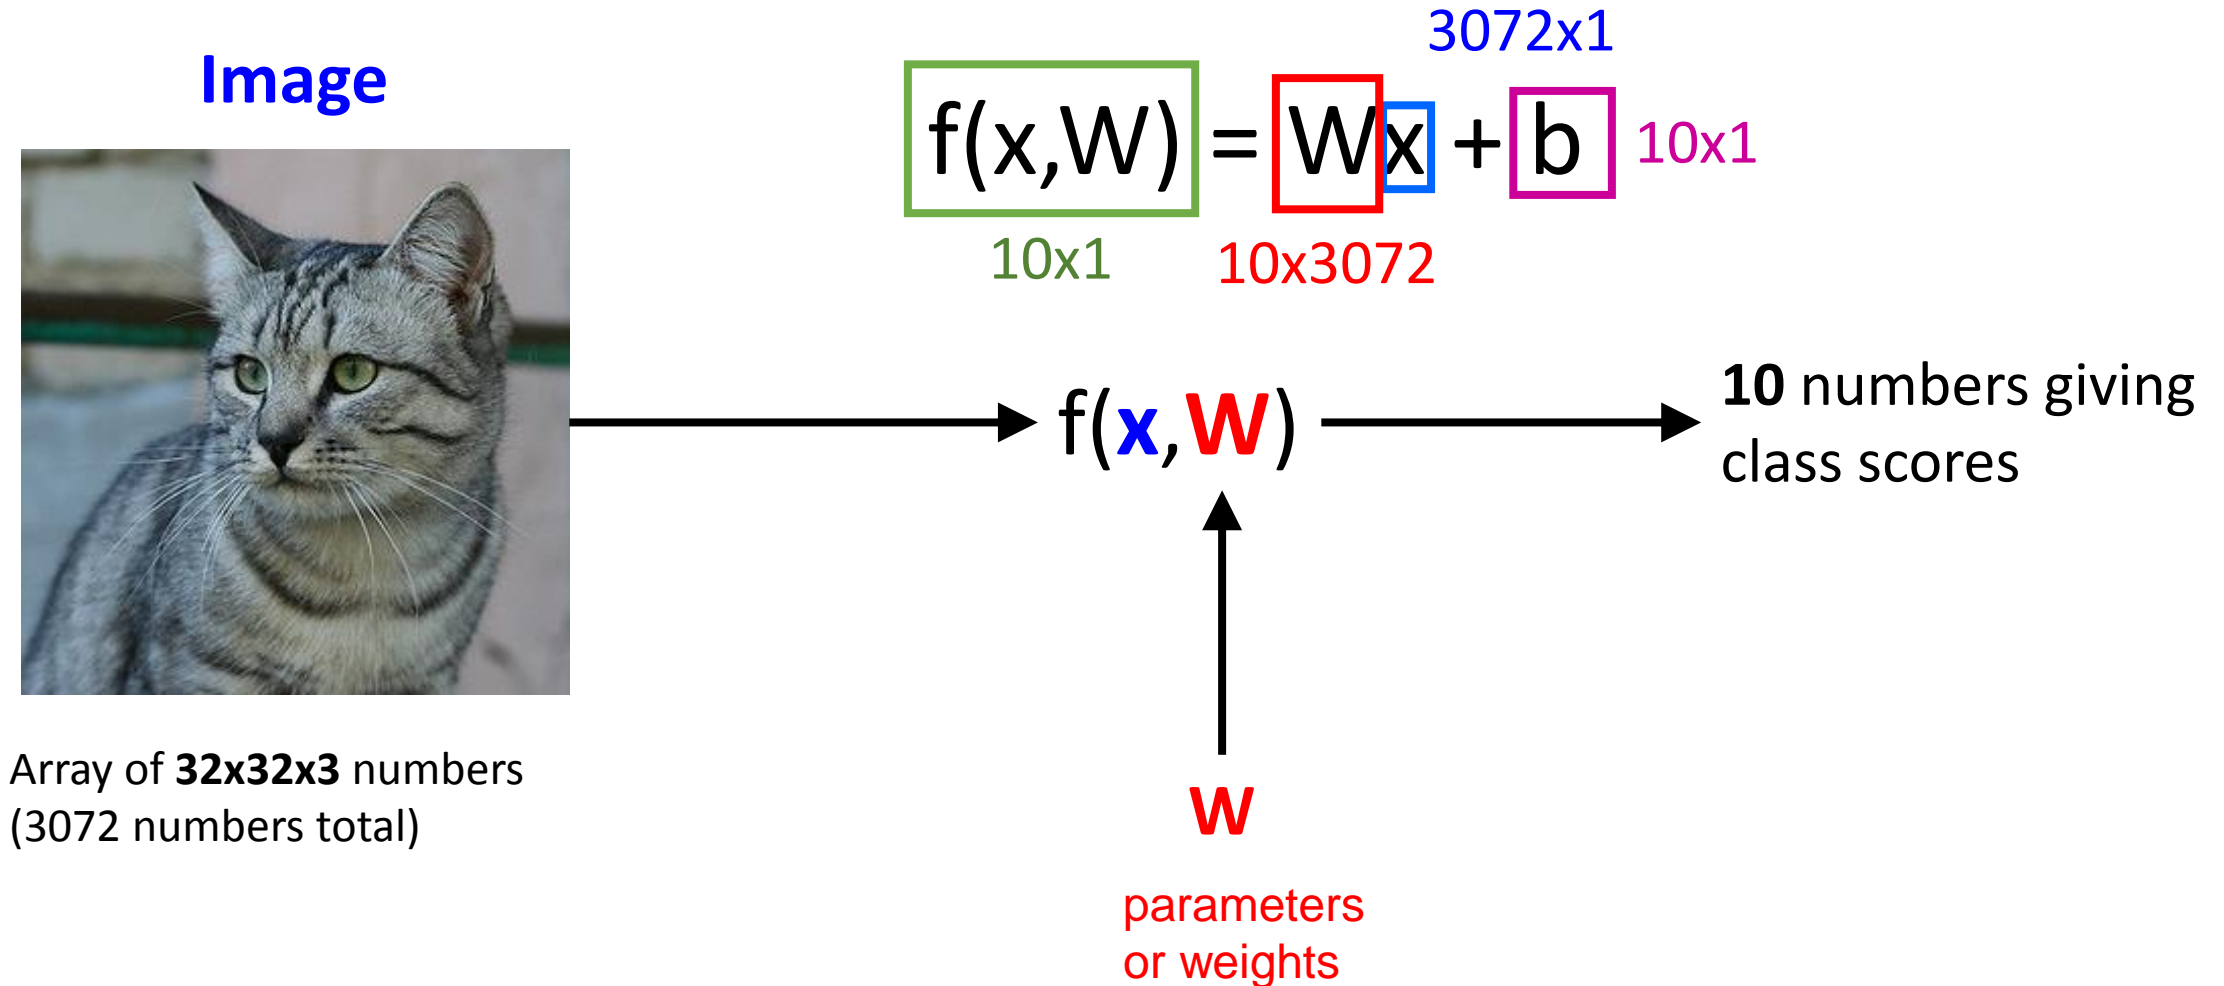
parameters  
or weights

# Parametric Approach: Linear Classifier



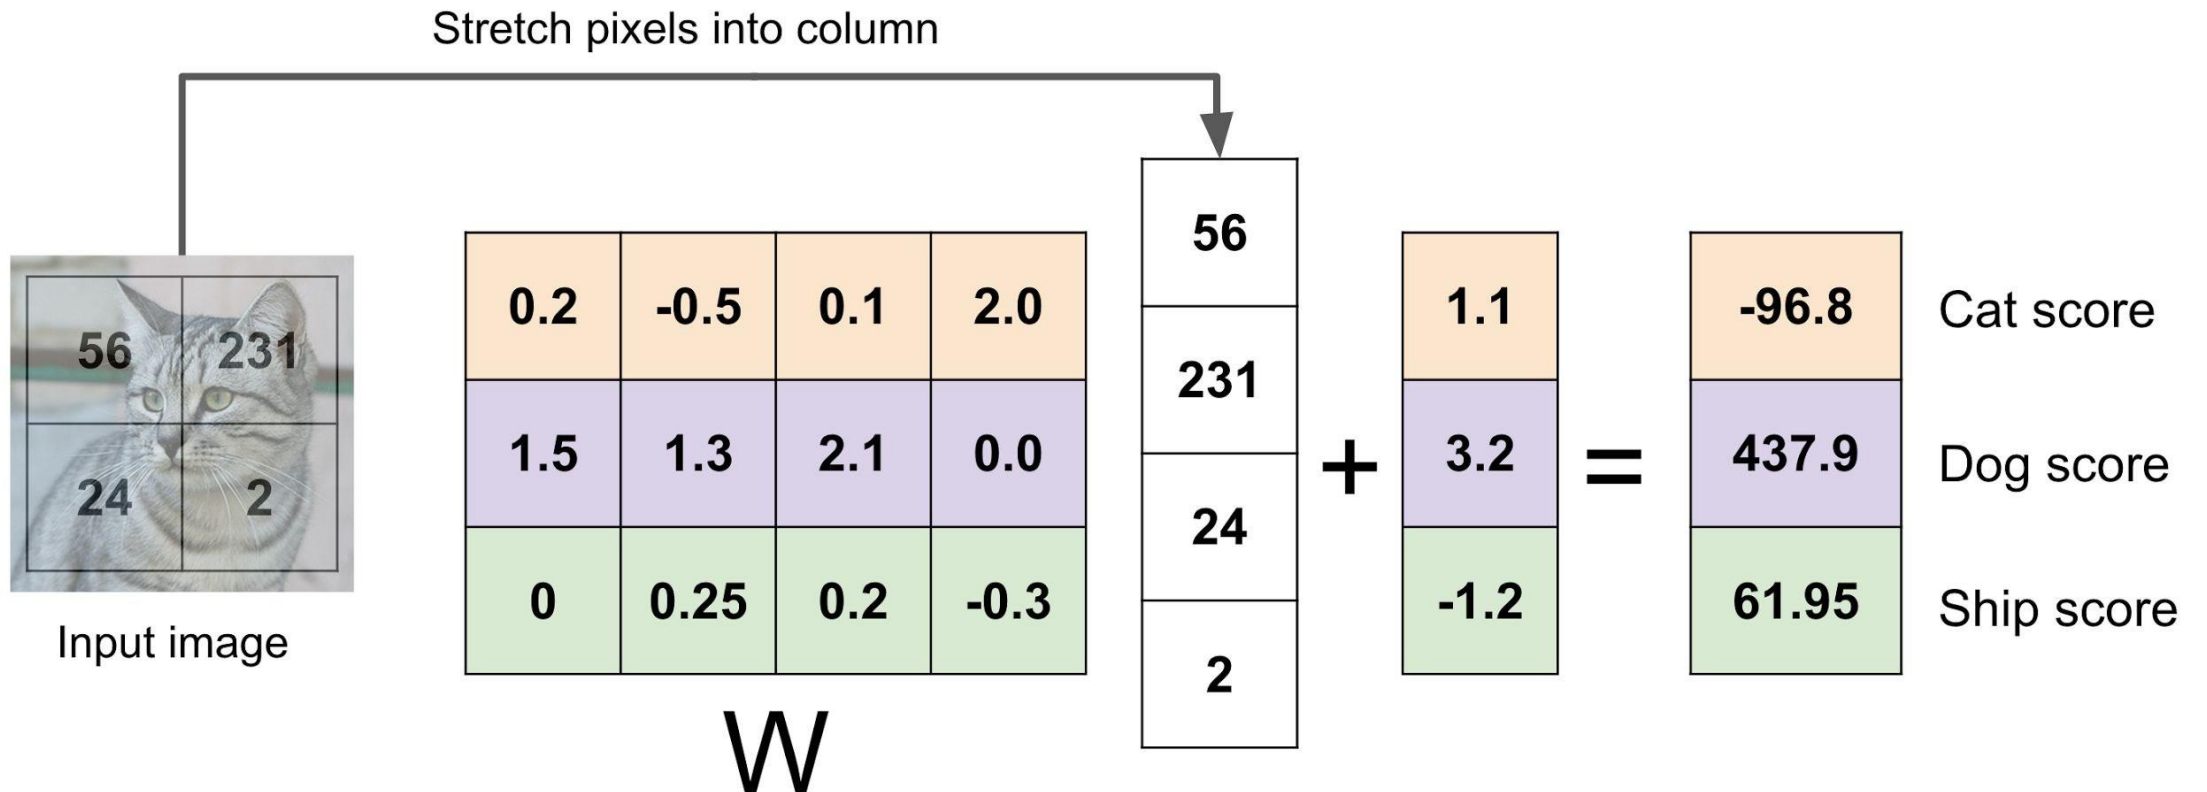


# Parametric Approach: Linear Classifier

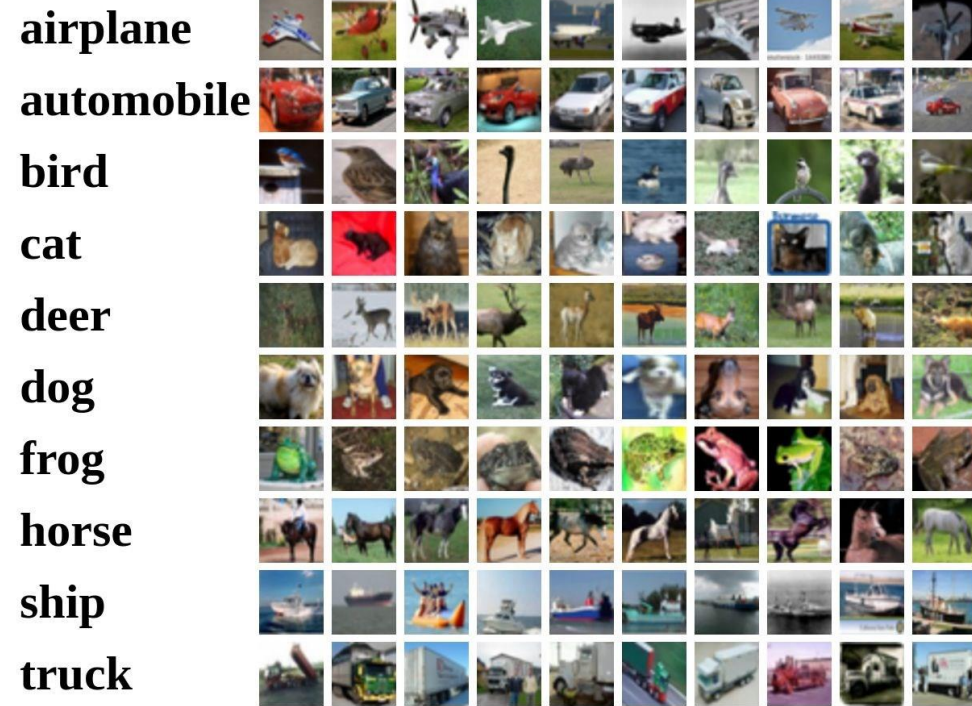


# Toy example

Image with 4 pixels, and 3 classes (cat/dog/ship)



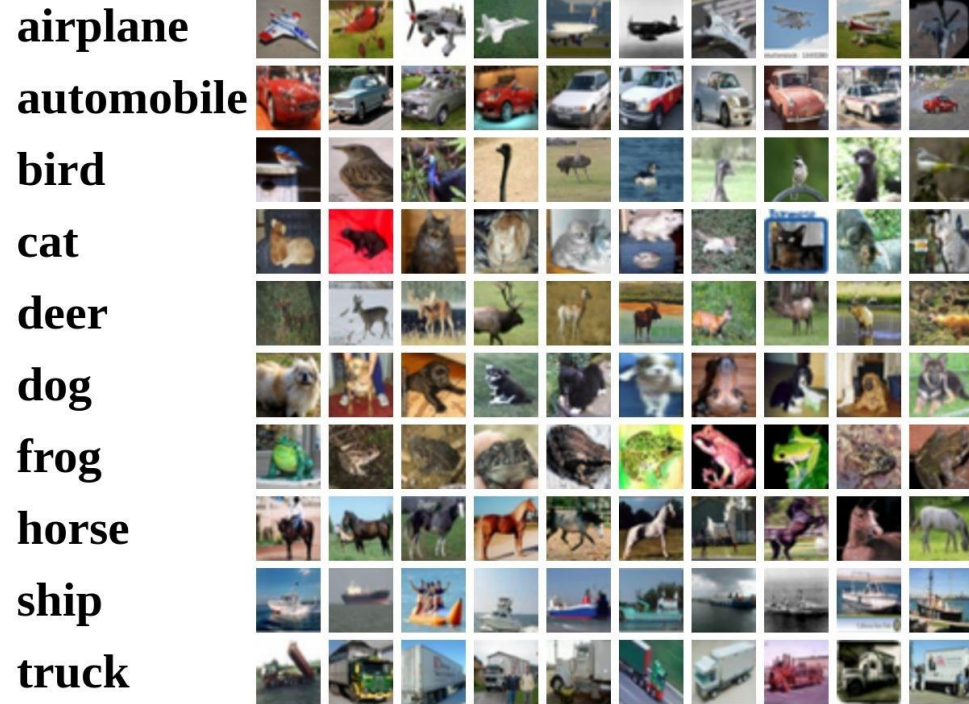
# Interpreting a Linear Classifier



$$f(x, W) = Wx + b$$

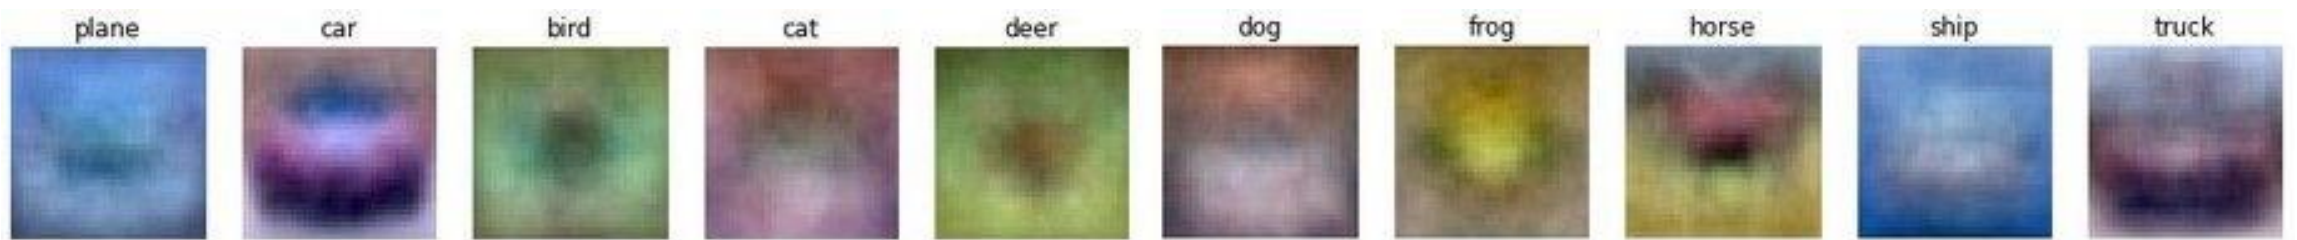
What is this thing doing?

# Interpreting a Linear Classifier

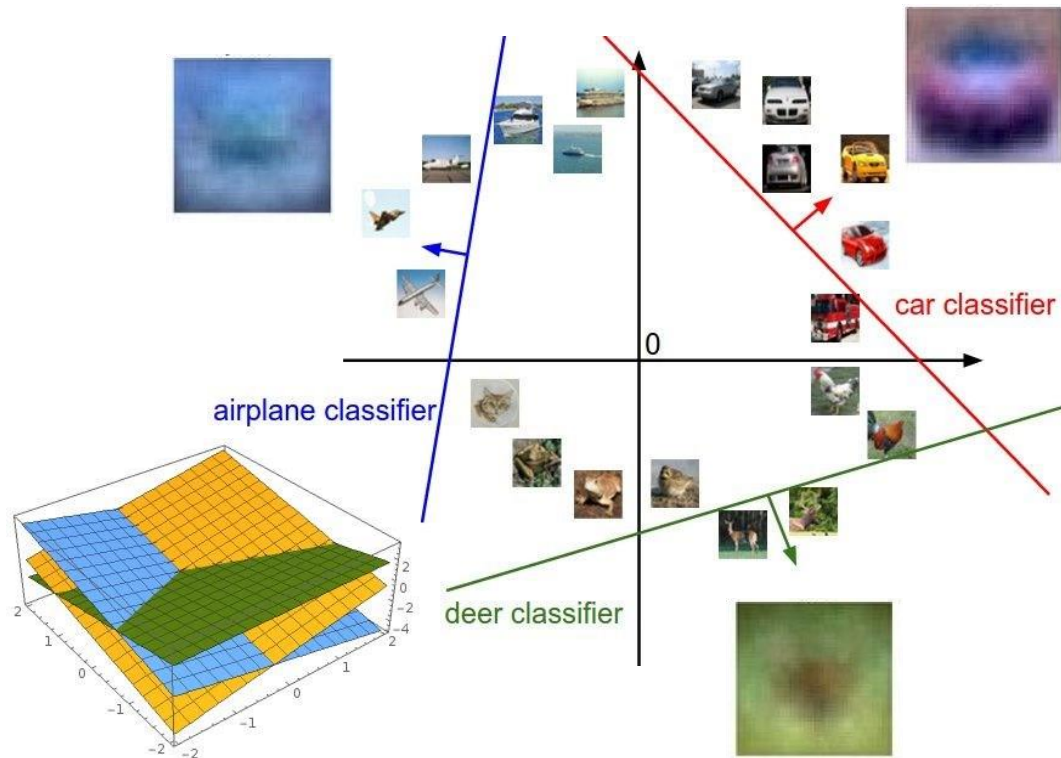


$$f(x, W) = Wx + b$$

Example trained weights  
of a linear classifier  
trained on CIFAR-10:



# Interpreting a Linear Classifier



$$f(x, W) = Wx + b$$



Array of **32x32x3** numbers  
(3072 numbers total)



# How to define level of unhappiness with the scores?



airplane	-3.45	-0.51	3.42
automobile	-8.87	<b>6.04</b>	4.64
bird	0.09	5.31	2.65
cat	<b>2.9</b>	-4.22	5.1
deer	4.48	-4.19	2.64
dog	8.02	3.58	5.55
frog	3.78	4.49	<b>-4.34</b>
horse	1.06	-4.37	-1.5
ship	-0.36	-2.09	-4.79
truck	-0.72	-2.93	6.14

# Toy example

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>

A **loss function** tells how good our classifier

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

Loss over the dataset is a sum of loss over examples:

$$L = \frac{1}{N} \sum_i L_i(f(x_i, W), y_i)$$

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$\begin{aligned} L_i &= \sum_{j \neq y_i} \begin{cases} 0 & \text{if } s_{y_i} \geq s_j + 1 \\ s_j - s_{y_i} + 1 & \text{otherwise} \end{cases} \\ &= \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1) \end{aligned}$$

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
Loss	<b>2.9</b>		

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

$$\begin{aligned} &= \max(0, 5.1 - 3.2 + 1) \\ &+ \max(0, -1.7 - 3.2 + 1) \\ &= \max(0, 2.9) + \max(0, -3.9) \\ &= 2.9 + 0 \\ &= 2.9 \end{aligned}$$

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
Loss	2.9	0	

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

$$= \max(0, 1.3 - 4.9 + 1)$$

$$+ \max(0, 2.0 - 4.9 + 1)$$

$$= \max(0, -2.6) + \max(0, -1.9)$$

$$= 0 + 0$$

$$= 0$$



Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
<b>Loss</b>	2.9	0	<b>12.9</b>

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

$$\begin{aligned}
 &= \max(0, 2.2 - (-3.1) + 1) \\
 &+ \max(0, 2.5 - (-3.1) + 1) \\
 &= \max(0, 6.3) + \max(0, 6.6) \\
 &= 6.3 + 6.6 \\
 &= 12.9
 \end{aligned}$$

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
Loss	2.9	0	12.9

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

Loss over the dataset  $L = \frac{1}{N} \sum_{i=1}^N L_i$

$$L = (2.9 + 0 + 12.9)/3 \\ = 5.27$$

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
Loss	2.9	0	12.9

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

Q: What happens to  
loss if car scores  
change a bit?

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
Loss	2.9	0	12.9

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

Q2: what is the  
min/max possible  
loss?

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
Loss	2.9	0	12.9

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

Q3: At initialization  $W$   
is small so all  $s \approx 0$ .  
What is the loss?



Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
Loss	2.9	0	12.9

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

Q4: What if the sum  
was over all classes?  
(including  $j = y_i$ )

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
Loss	2.9	0	12.9

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

Q5: What if we used  
mean instead of  
sum?

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
<b>Loss</b>	<b>2.9</b>	<b>0</b>	<b>12.9</b>

## Multiclass SVM (hinge) loss:

$x_i$  - image

$y_i$  - label, element of a set  $\{0, 1, \dots\}$

scores vector

$$s = f(x_i, W) = [s_0, \dots, s_{y_i}, \dots]$$

$s_{y_i}$  element corresponds  
to ground truth label  $y_i$

SVM loss

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

Q6: What if we used

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)^2$$

$$f(x, W) = Wx$$

$$L = \frac{1}{N} \sum_{i=1}^N \sum_{j \neq y_i} \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + 1)$$

E.g. Suppose that we found a  $W$  such that  $L = 0$ .  
Is this  $W$  unique?

**No!  $2W$  is also has  $L = 0$ !**

Consider: 3 training examples, 3 classes.  
With the scores



cat	<b>3.2</b>	1.3	2.2
car	5.1	<b>4.9</b>	2.5
frog	-1.7	2.0	<b>-3.1</b>
Loss	2.9	0	

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

Before:

$$\begin{aligned} &= \max(0, 1.3 - 4.9 + 1) \\ &+ \max(0, 2.0 - 4.9 + 1) \\ &= \max(0, -2.6) + \max(0, -1.9) \\ &= 0 + 0 \\ &= 0 \end{aligned}$$

With 2W:

$$\begin{aligned} &= \max(0, 2.6 - 9.8 + 1) \\ &+ \max(0, 4.0 - 9.8 + 1) \\ &= \max(0, -6.2) + \max(0, -4.8) \\ &= 0 + 0 \end{aligned}$$



# Regularization

$$L = \underbrace{\frac{1}{N} \sum_{i=1}^N \sum_{j \neq y_i} \max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + 1)}_{\text{Data loss}} + \underbrace{\lambda R(W)}_{\text{Regularization}}$$

$\lambda$  - regularization strength (hyperparameter)

**L2 regularization**

$$R(W) = \sum_k \sum_l W_{k,l}^2$$

**L1 regularization**

$$R(W) = \sum_k \sum_l |W_{k,l}|$$

**Elastic net (L1 + L2)**

$$R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$$

**Dropout (will see later)**

**Batch normalization (will see later)**

# L2 regularization: motivation

$$x = [1, 1, 1, 1] \qquad R(W) = \sum_k \sum_l W_{k,l}^2$$

$$w_1 = [1, 0, 0, 0]$$

$$w_2 = [0.25, 0.25, 0.25, 0.25]$$

$$w_1^T x = w_2^T x = 1$$

# Softmax Classifier (Multinomial Logistic Regression)

scores = unnormalized log probabilities of the classes.

$$s = f(x_i, W)$$

$$P(Y = y_i | X = x_i) = \frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \quad \text{probability of } x_i \text{ image has } y_i \text{ label}$$



cat	<b>3.2</b>
car	5.1
frog	-1.7

# Softmax Classifier (Multinomial Logistic Regression)

scores = unnormalized log probabilities of the classes.

$$s = f(x_i, W)$$

$$P(Y = y_i | X = x_i) = \frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \text{ probability of } x_i \text{ image has } y_i \text{ label}$$



cat	<b>3.2</b>
car	5.1
frog	-1.7

Softmax function

# Softmax Classifier (Multinomial Logistic Regression)



cat	<b>3.2</b>
car	5.1
frog	-1.7

scores = unnormalized log probabilities of the classes.

$$s = f(x_i, W)$$

$$P(Y = y_i | X = x_i) = \frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \quad \text{probability of } x_i \text{ image has } y_i \text{ label}$$

Want to maximize the likelihood  $\prod_i P(Y = y_i | X = x_i)$

or log likelihood  $\sum_i \log P(Y = y_i | X = x_i)$

or to minimize the negative log likelihood of the correct class:

$$L_i = -\log P(Y = y_i | X = x_i)$$



# Softmax Classifier (Multinomial Logistic Regression)



cat	<b>3.2</b>
car	5.1
frog	-1.7

scores = unnormalized log probabilities of the classes.

$$s = f(x_i, W)$$

$$P(Y = y_i | X = x_i) = \frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \quad \text{probability of } x_i \text{ image has } y_i \text{ label}$$

Want to minimize the negative log likelihood of the correct class:

$$L_i = -\log P(Y = y_i | X = x_i)$$

or

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$

# Softmax Classifier (Multinomial Logistic Regression)

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$



$$y_i = 0$$

cat	<b>3.2</b>
car	5.1
frog	-1.7

unnormalized log probabilities

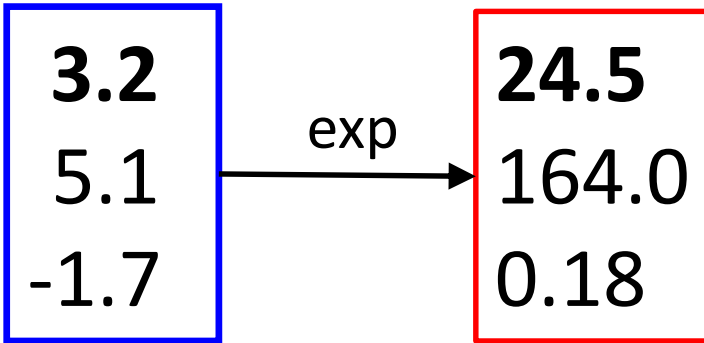
# Softmax Classifier (Multinomial Logistic Regression)

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$



unnormalized probabilities

cat  
car  
frog



unnormalized log probabilities

# Softmax Classifier (Multinomial Logistic Regression)

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$



unnormalized probabilities

cat  
car  
frog

**3.2**  
5.1  
-1.7

exp

**24.5**  
164.0  
0.18

normalize

**0.13**  
0.87  
0.00

unnormalized log probabilities

probabilities

# Softmax Classifier (Multinomial Logistic Regression)

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$



unnormalized probabilities

cat  
car  
frog

<b>3.2</b>
5.1
-1.7

exp

<b>24.5</b>
164.0
0.18

normalize

<b>0.13</b>
0.87
0.00

$y_i = 0$

$$L_i = -\log(0.13) = 0.89$$

unnormalized log probabilities

probabilities



# Softmax Classifier (Multinomial Logistic Regression)

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$



unnormalized probabilities

cat  
car  
frog

<b>3.2</b>
5.1
-1.7

exp

<b>24.5</b>
164.0
0.18

normalize

<b>0.13</b>
0.87
0.00

$y_i = 0$

$$L_i = -\log(0.13) = 0.89$$

unnormalized log probabilities

probabilities

# Softmax Classifier (Multinomial Logistic Regression)



unnormalized probabilities

cat  
car  
frog

**3.2**  
5.1  
-1.7

exp

**24.5**  
164.0  
0.18

normalize

**0.13**  
0.87  
0.00

unnormalized log probabilities

probabilities

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$

Q: What is the possible min/max loss  $L_i$ ?

$$y_i = 0$$

$$L_i = -\log(0.13) = 0.89$$

# Softmax Classifier (Multinomial Logistic Regression)



unnormalized probabilities

cat  
car  
frog

**3.2**  
5.1  
-1.7

exp

**24.5**  
164.0  
0.18

normalize

**0.13**  
0.87  
0.001

unnormalized log probabilities

probabilities

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$

Q2: Usually at initialization  $W$  is small so all  $s \approx 0$ .  
What is the loss?

$$y_i = 0$$

$$L_i = -\log(0.13) = 0.89$$

# Cross-entropy loss (softmax loss)

Kullback–Leibler divergence – **nonsymmetrical**  
distance between two probability distribution  
 $P$  and  $Q$

$$D_{KL}(P||Q) = \sum_x p(x) \log \frac{p(x)}{q(x)} = \underbrace{\sum_x p(x) \log p(x)}_{H(P) \text{ entropy of } P} - \underbrace{\sum_x p(x) \log q(x)}_{H(P, Q) \text{ cross-entropy of } P \text{ and } Q} = H(P) + H(P, Q)$$

Image classification:

$P(Y = y_i | X = x_i)$  – “true” distribution of class probabilities

$Q(Y = \hat{y}_i | X = x_i)$  – estimated class probabilities, softmax output of classifier

# Cross-entropy loss (softmax loss)

Want to minimize Kullback–Leibler divergence

$$D_{KL}(P||Q) = \underbrace{\sum_i P(y_i|x_i) \log P(y_i|x_i)}_{\text{constant}} - \underbrace{\sum_i P(y_i|x_i) \log Q(\hat{y}_i|x_i)}_{\mathbf{H(P, Q)} \text{ cross-entropy of } P \text{ and } Q}$$

minimizing Kullback–Leibler divergence is equivalent to  
minimizing cross-entropy of  $P$  and  $Q$



# Cross-entropy loss (softmax loss)

Want to minimize cross entropy loss for an image  $x_i$

$$L_i = -P(y_i|x_i)\log Q(\hat{y}_i|x_i)$$

If image labeled with single class  $y_i$  then

one-hot label encoding

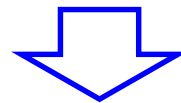
$$P(y_i|x_i) = [0, \dots 1, \dots 0]$$



single 1 at  $y_i$ -th position

softmax of classifier scores

$$Q(\hat{y}_i|x_i) = \frac{e^{s_i}}{\sum_j e^{s_j}}$$



$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$

# Cross-entropy loss (softmax loss)

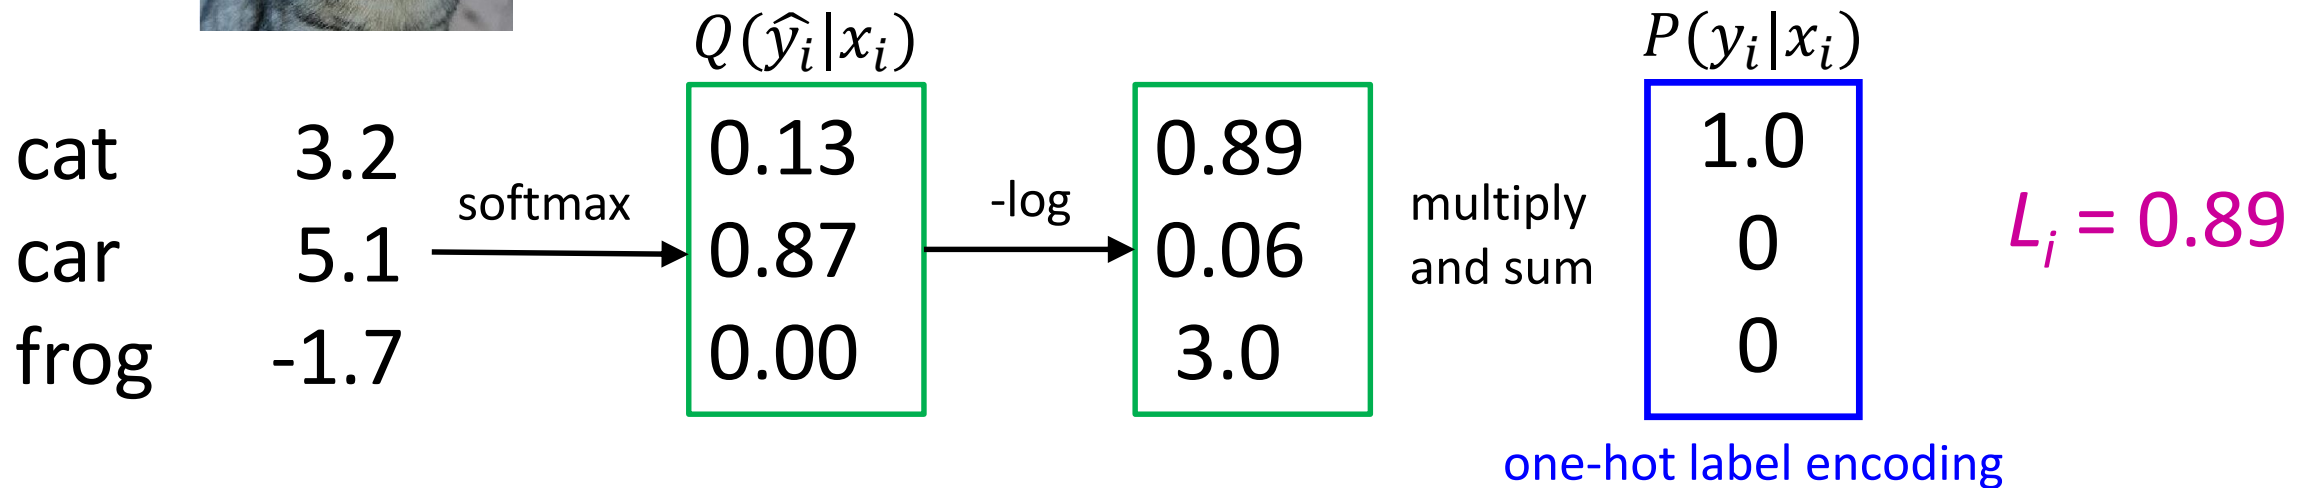
Want to minimize cross entropy loss for an image  $x_i$

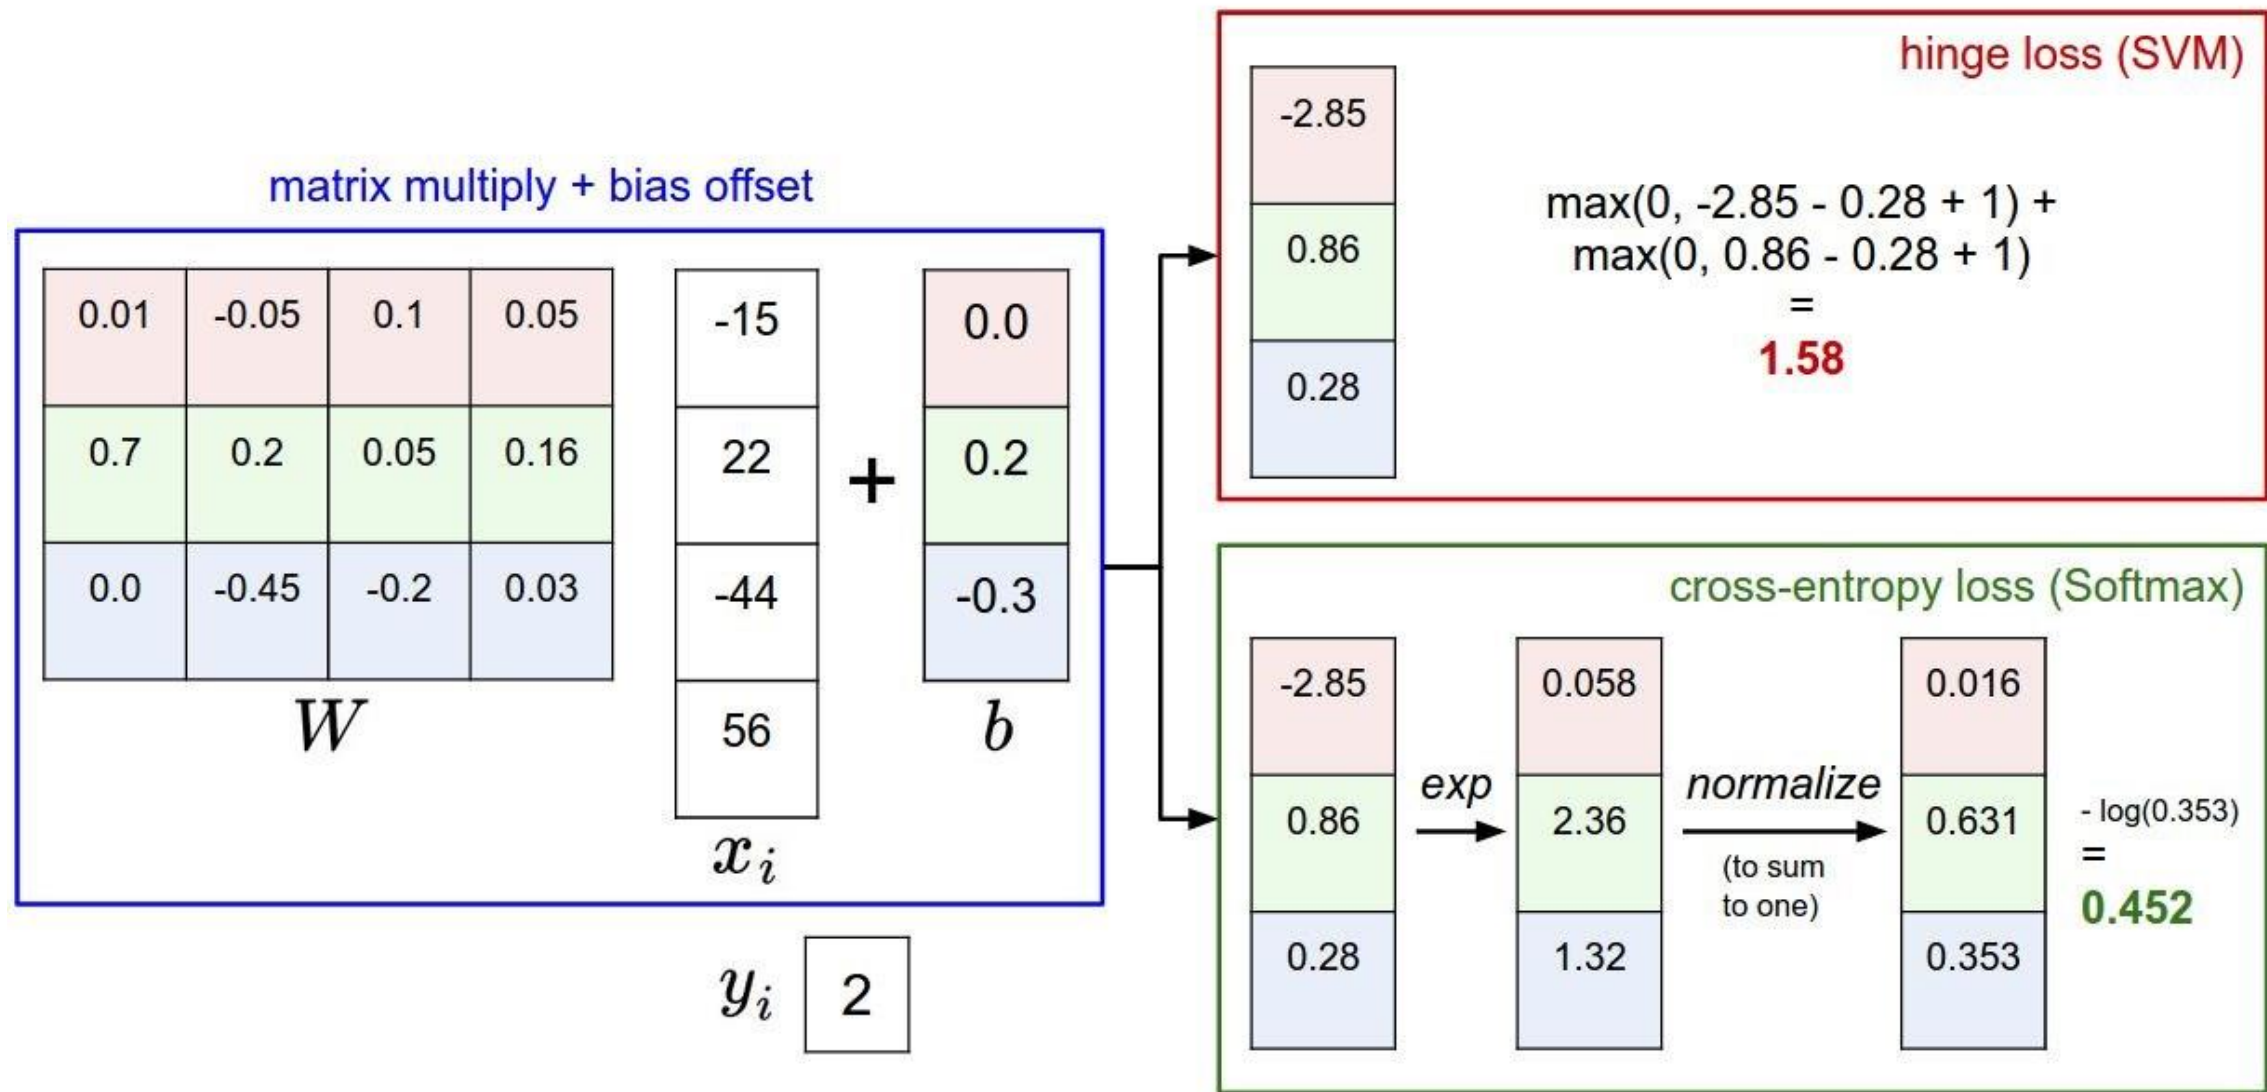
$$L_i = -P(y_i|x_i)\log Q(\hat{y}_i|x_i)$$



$P(y_i|x_i)$  - ground truth probabilities

$Q(\hat{y}_i|x_i) = \frac{e^{s_i}}{\sum_j e^{s_j}}$  - softmax of classifier scores





# Softmax vs SVM

Softmax

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$

VS

SVM

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$



# Softmax vs SVM

Softmax

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}}$$

vs

SVM

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1)$$

assume scores:

[10, -2, 3]

[10, 9, 9]

[10, -100, -100]

and  $y_i = 0$

Q: Suppose I take a datapoint and I jiggle a bit (changing its score slightly). What happens to the loss in both cases?



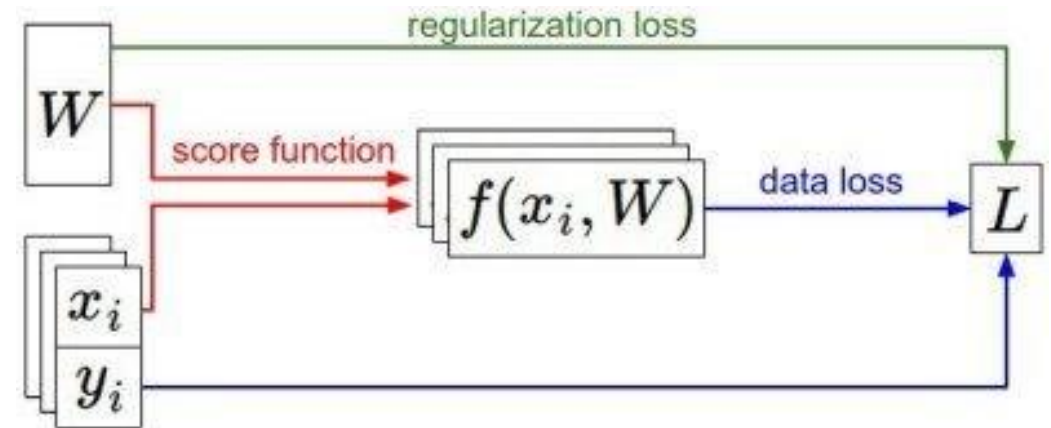
# Recap

- We have some dataset of  $(x_i, y_i)$
- We have a **score function**:  $s = f(x_i, W)$
- We have a **loss function**:

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \quad \text{Softmax}$$

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1) \quad \text{SVM}$$

$$L = \frac{1}{N} \sum_{i=1}^N L_i + R(W) \quad \text{Full loss}$$



# Recap

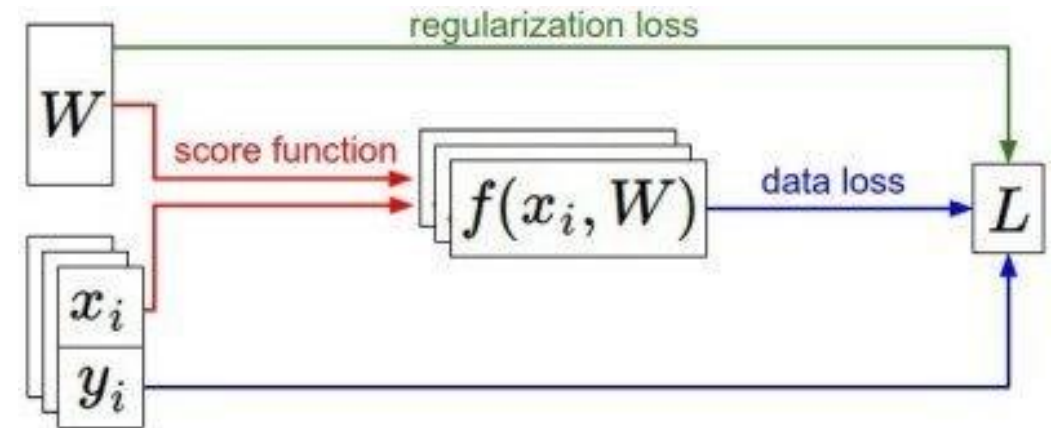
How do we find the best  $W$ ?

- We have some dataset of  $(x_i, y_i)$
- We have a **score function**:  $s = f(x_i, W)$
- We have a **loss function**:

$$L_i = -\log \frac{e^{s_{y_i}}}{\sum_j e^{s_j}} \quad \text{Softmax}$$

$$L_i = \sum_{j \neq y_i} \max(0, s_j - s_{y_i} + 1) \quad \text{SVM}$$

$$L = \frac{1}{N} \sum_{i=1}^N L_i + R(W) \quad \text{Full loss}$$



# Next time

Optimization

Introduction to neural networks

Backpropagation